

**Midterm Exam, Operating Systems, Spring 2022,
Department of CSIE, NTUT, May 2, 2022**

Note: There are 8 questions worth a total of 110 points. You have 110 minutes. The exam is ***open book***. Any form of discussions are ***strictly prohibited*** during the exam.

1. **(10pt)** Modern computers are designed based on *von Neumann architecture*.
 - (1) Please list the major components of a computer system, and describe how they are connected. (5pt)
 - (2) Please give the detailed steps of actions taken by these components during an I/O operation (e.g. reading a file). (5pt)

2. **(10pt)** What's the purpose of *interrupts*? How do we usually handle interrupts in OS kernel? Why are modern operating systems usually interrupt-driven? What are the benefits of this design?

3. **(20pt)** Regarding the following questions about *process and thread management*, please indicate whether each statement is *true* or *false*. **If a statement is incorrect, please explain the reasons to get the full score.** (*not* just correcting the error)
 - (1) Kernel-level threads are directly supported and managed by OS kernel. (4pt)
 - (2) Multiple user-level threads, as managed by programmers, can always be run in parallel on multicore systems. (4pt)
 - (3) The dominant factor of performance gains from adding additional CPU cores to an application is the number of cores. (4pt)
 - (4) Multithreaded programs can always provide better performance than a single-threaded solution. (4pt)
 - (5) A multithreaded program using multiple user-level threads achieves better performance on a multiprocessor system than on a single-processor system. (4pt)

4. **(10pt)** Answer the following questions regarding *deadlocks*:
- (1) Please explain the necessary conditions for deadlocks. (5pt)
 - (2) Among the different methods for handling deadlocks, what are the differences between *deadlock prevention* and *deadlock avoidance*? (5pt)
5. **(10pt)** Consider the following snapshot of a system with five processes P_0 through P_4 and four resource types A, B, C, and D:

	<u>Allocation</u>	<u>Max</u>	<u>Available</u>
	A B C D	A B C D	A B C D
P_0	3 0 1 3	5 1 1 8	1 0 0 2
P_1	2 2 4 0	3 2 4 1	
P_2	3 1 2 1	3 4 2 1	
P_3	1 4 0 1	4 6 3 3	
P_4	2 2 1 2	6 3 2 5	

Answer the following questions using the *banker's algorithm*:

- (1) Determine whether or not the system is safe. If the state is safe, illustrate the order in which the processes may complete. Otherwise, illustrate why the state is unsafe. (5pt)
 - (2) If a request from process P_4 arrives for (0, 0, 0, 2), can the request be granted immediately? Please justify your answer. (5pt)
6. **(15pt)** Answer the following questions regarding memory management.
- (1) What are the differences between *pure paging* and *pure segmentation* in terms of the following aspects: fragmentation, address translation structures, ability to share code across processes, and memory protection? (10pt)
 - (2) Consider a paging system with the page table stored in memory. Assume that a memory reference takes 80 nanoseconds. If we add TLBs, and we want the effective memory access time to be less than 10% slowdown, what is the minimum percentage of cache hit in the TLBs? (Assume that finding a page-table entry in the TLBs takes 4 nanoseconds, if the entry is present.) (5pt)

7. (15pt) Answer the following questions regarding synchronization. The simplest software tool to solve the critical-section problem is the *mutex lock*. Instead of implementing mutex locks using atomic hardware instructions, we can simply implement two operations for acquiring and releasing the lock as follows:

```

acquire()
{
    while (!available)
        ;
    available = false;
}
release()
{
    available = true;
}

```

- (1) What's the problem with this implementation? Please compare the effects of such implementation on uniprocessor and multiprocessor systems, respectively. (5pt)
- (2) What changes would be necessary so that a process waiting to acquire a mutex lock would be blocked and placed into a waiting queue until the lock became available? (5pt)
- (3) What is the idea of *semaphore*? Please compare its difference in usage with mutex locks. (5pt)

8. (20pt) Consider the following set of processes, with the length of the CPU burst given in milliseconds:

Process	Priority	Burst Time
P ₁	2	3
P ₂	1	2
P ₃	4	10
P ₄	2	6
P ₅	3	4

The processes are assumed to have arrived in the order P₁, P₂, P₃, P₄, P₅, all at time 0. Note that a *small* priority number implies a *high* priority. Answer the questions using the following scheduling algorithms: **SJF**, **non-preemptive priority**, and **RR** (time quantum=2).

- (1) What is the *turnaround time* of each process for each of the scheduling algorithms? (10pt)
- (2) What is the *waiting time* of each process for each of the scheduling algorithms? Which of the algorithms results in the minimum average waiting time (over all processes)? (10pt)